

$P \Rightarrow Q \equiv \neg P \vee Q$

 $\{t / (\exists u) (u \notin \text{Branch} \vee u[2] \neq 'KGIP') \vee (\exists v) (v \in \text{Deposit} \wedge u[1] = v[1] \wedge t[1] = v[3])\}$

Midsem syllabus: Tuple Calculus

* SQL
 select A_1, A_2, \dots, A_n from R_1, R_2, \dots, R_m where P ;
 $\equiv \Pi_{A_1, A_2, \dots, A_n} (\sigma_p (R_1 \times R_2 \times \dots \times R_m))$

union intersect minus
and or not

(i) Find the name of all branches in Deposit relation.

select b-name
from Deposit;
Select distinct (b-name)
from Deposit;

(ii) Find all customers who have account at IIT Branch.

select c-name

from Deposit

where b-name = 'IIT';

(iii) Find all customers having a loan or account or both at IIT Branch.

Select c-name

Select c-name
from Deposit
where b-name = 'IIT' } union {
select c-name
from Deposit
where b-name = 'IIT'

(iv) Find all customers having a loan at any branch
and their city.

Select c.c-name, c-city

from Customer C, Borrower B

where C.c-name = B.c-name

and b-name = 'IIT'

(v) Find all customers who have both loan and
account at IIT Branch.

Select c-name

from ~~Customer~~ Deposit

where b-name = 'IIT'

and ~~c-name~~ c-name in

(Select c-name)

from Borrower B

where b-name = 'IIT'

(vi) Find all customers who have account at IIT Branch
but no loan from IIT Branch.

Select c-name

from Deposit

where b-name = 'IIT'

(Select c-name
from Borrow
where b-name = 'IIT')

(vii) Find all branches that have more assets than any branch located in KGP.

Select T.b-name
from Branch T, Branch S
where T.asset > S.asset
and S.b-city = 'KGP'

Select b-name
from Branch
where asset > any

(Select asset

from Branch
where b-city = 'KGP')

> any < any > all < all

(viii) Find all customers with balance lies between 50000 and 60000.

Select c-name
from Deposit
where Balance > 50000 and Balance < 60000

Select c-name
from Deposit
where Balance between
50000 and 60000

(ix) List in alphabetical order all the customers having a loan at IIT Branch.

Select c-name

from Borrow

where b-name = 'IIT'

order by c-name,

avg, min, max, sum, count

(x) Find the average account balance at all branches.

Select b-name, avg(balance)

from Deposit

group by b-name.

(xi) Find the average account balance at all branches having average balance greater than 10000.

Select b-name, avg(balance)

from deposit

group by b-name

having avg(balance) > 10000

(xii) Find no. of tuples in customer relation.

Select count(*)

from Customer

(xiii) Find all customers who have deposit at IIT Branch but for whom there is no entry in the customer relation.

Select c-name
from Deposit
where b-name = 'IIT'
and c-name not in

Select c-name
from Deposit
where b-name = 'IIT'
and not exists

Select *
from customer C
where D.c-name = C.c-name

Example : Teacher (T-name, Dept, Tel-no, S-title)

Student (S-name, ~~course~~, hall)

Study (S-title, S-name, level, status, marks)

Find the name of teachers of Maths department who teach 200 level subjects.

Select T-name
from Teacher, Study
where Dept = 'Maths'
and level = '200'
and Teacher.S-title = Study.S-title

Find students of CSDP course living in VS or SN hall who study no subject taught by professor XYZ.

Select s-name
from student
where course = 'CSDP' and
(hall = 'VS' or hall = 'SN')
and s-name not in
(select s-name
from Teacher T, Study S
where t-name = 'XYZ' and
and T.s-title = S.s-title)

OR
SELECT s-name
FROM STUDENT S
WHERE S-COURSE = 'CSDP'
AND S-HALL IN ('VS', 'SN')
AND S-NOT IN
(SELECT S-NOME
FROM TEACHER T,
STUDY S
WHERE T-TNAME = 'XYZ'
AND T-S-TITLE = S-S-TITLE)

SELECT s-name
FROM STUDENT S
WHERE S-COURSE = 'CSDP'
AND S-HALL IN ('VS', 'SN')
AND S-NOT IN
(SELECT S-NOME
FROM TEACHER T,
STUDY S
WHERE T-TNAME = 'XYZ'
AND T-S-TITLE = S-S-TITLE)

SELECT s-name
FROM STUDENT S
WHERE S-COURSE = 'CSDP'
AND S-HALL IN ('VS', 'SN')
AND S-NOT IN
(SELECT S-NOME
FROM TEACHER T,
STUDY S
WHERE T-TNAME = 'XYZ'
AND T-S-TITLE = S-S-TITLE)

SELECT s-name
FROM STUDENT S
WHERE S-COURSE = 'CSDP'
AND S-HALL IN ('VS', 'SN')
AND S-NOT IN
(SELECT S-NOME
FROM TEACHER T,
STUDY S
WHERE T-TNAME = 'XYZ'
AND T-S-TITLE = S-S-TITLE)

SELECT s-name
FROM STUDENT S
WHERE S-COURSE = 'CSDP'
AND S-HALL IN ('VS', 'SN')
AND S-NOT IN
(SELECT S-NOME
FROM TEACHER T,
STUDY S
WHERE T-TNAME = 'XYZ'
AND T-S-TITLE = S-S-TITLE)

QWEL

range of t_1 is R_1

range of t_2 is R_2

range of t_n is R_n

retrieve $(t_1, A_j, t_2, A_{j_2}, \dots, t_m, A_{j_m})$

where P_i

Q. Find all customers having an account at IIT Branch.

range of t is Deposit

retrieve $(t.c_name)$

where $t.b_name = 'IIT'$

Q. Find all customers having a loan at IIT Branch and their cities.

range of t is Borrow

range of u is Customer

retrieve $(t.c_name, u.c_city)$

where $t.c_name = u.c_name$

and $t.b_name = 'IIT'$

8. find all customers who have loan and account ~~is at~~
at IIT Branch.
- range of t is ~~Borrow~~ Borrow,
range of u is Deposit,
retrieve (t.c-name)
where t.c-name = u.c-name and
t.b-name = 'IIT' and
u.b-name = 'IIT'
- use same logic as file 7 to combine the above
- g. find all ~~is~~ customers who have account or loan
or both at IIT Branch.
- range of u is Deposit
retrieve into temp (u.c-name)
where u.b-name = 'IIT'
range of s is Borrow
append to temp (s.c-name)
where s.b-name = 'IIT'
- Pointing:
- range of t is temp,
retrieve unique (t.c-name)
- Q. find all customers who have account at IIT Branch
but dont have loan at IIT Branch.
- range of u is Deposit
retrieve into temp (u.c-name)

range of t is temp
delete (t)

where $t.c_name = s.c_name$ and
 $s.b_name = 'IIT'$

Pointing:

range of t is temp
retrieve unique ($t.c_name$)

Q. Find all customer names in alphabetical order who have account at IIT Branch.

range of u is Deposit

retrieve $u.c_name$

where $u.b_name = 'IIT'$

sort by $u.c_name$

max min avg sum count \rightarrow aggregate operators

Q. Find the average account balance for all accounts at 'IIT' Branch.

range of t is Deposit

retrieve avg ($t.balance$, where $t.b_name = 'IIT'$)

Q. Find all account numbers whose balance is more than the average balance at the branch where the account is held.

range of t is Deposit
retrieve ($t \cdot \text{acc_no}$) is $t \cdot \text{balance} > \text{avg} (t \cdot \text{balance} \text{ by } t \cdot \text{b_name})$

Q. find all customers who have account at all branches in the city 'KGP'.

range of t is Deposit

range of u is Branch

range of s is Deposit

retrieve ($t \cdot c\text{-name}$)

where

countu ($s \cdot b\text{-name}$ where

$s \cdot b\text{-name} = u \cdot b\text{-name}$ and

$u \cdot b\text{-city} = 'KGP'$ and

$s \cdot c\text{-name} = t \cdot c\text{-name}$)

= countu ($s \cdot b\text{-name}$ #

countu ($u \cdot b\text{-name}$ #

where $u \cdot b\text{-city} = 'KGP'$)

Q. Find the name of subjects, such that all students studying the subject secure more than 80 marks.

range of t is study

retrieve (t , ~~$t.s_title$~~)

where $\min(t.\text{marks by } t.s_title) > 80$

Q. Find the students of CSDP, living in VS Hall or in SN Hall who study no subject taught by professor XYZ.

range of t is student

append to temp1 ($t.s_name$)

where ($t.course = 'CSDP'$ and ($t.hall = 'VS'$ or $t.hall = 'SN'$))

range of u is teacher

range of s is study

append to temp2 ($s.s_name$)

where $u.s_title = s.s_title$

and $u.t_name = 'XYZ'$.

Point :

range of $t1$ is temp1

range of $t2$ is temp2

delete($t1$)

where $t1.s_name = t2.s_name$.

range of t3 is tempt
retrieve unique (t3.s-name)

Q. find the maximum marks scored in each subject
of the CSDP course.

range of t is study

range of u is student

retrieve t.s-title, max (t.marks by
(t.s-title))

where t.s-name = u.s-name

and u.course = 'CSDP'

Q. find the name of students who secured equal
marks in 2 different subjects

range of t is study

range of u is study

retrieve (t.s-name)

where t.s-name \neq u.s-name

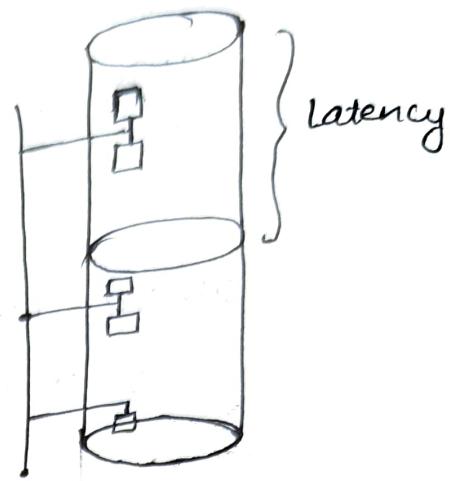
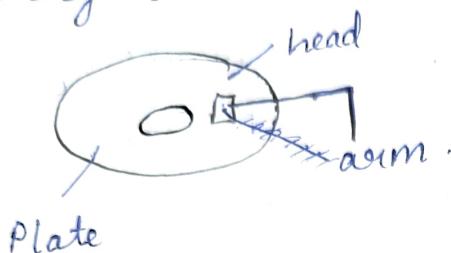
and t.s-title $< >$ u.s-title

and t.marks = u.marks.

FILE ORGANISATION

1 Block = 2400 Bytes

Data is transferred from primary to secondary memory block-wise.



Disk latency time.

Seek time.

seek time. It is the time it takes to move the arm to the correct cylinder.

Once the ~~arm~~ arm is placed on the correct cylinder, it takes some time to rotate and read the required data.

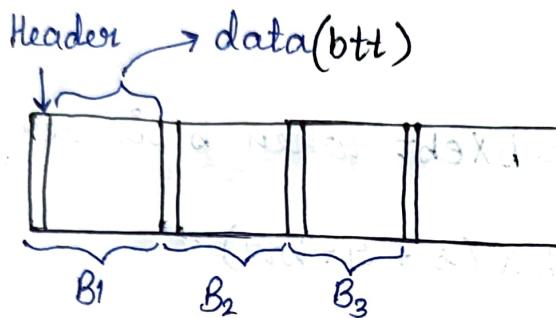
Average seek time(s) is the time it takes to ~~travel~~ ~~traverse~~ ~~1/3 rd~~ of the cylinder.

Average rotational latency time(r_1)

It is the average time it takes to locate a the correct block on the track once ~~the~~ the head is placed in the track. = $\frac{1}{2}$ revolution time

Block transfer time (btt)

btt is the time it takes to transfer the information of one block from three 2 memories once the head has been correctly placed at the beginning of the block.



Effective Block Transfer Time (EBT): Time taken to

$$\text{ebt} = \frac{B}{b}$$

transfer the information of one block along with the header.

Parameters for IBM 3980

B 2400 bytes

btt 0.8 ms

c 600

ebt 0.84 ms

N 885

τ_1 8.3 ms

δ 16 ms

Read 10 adjacent blocks

$$\text{time} = s + g_1 + 10 \times b \times t$$

$$= 32.7 \text{ ms.}$$

Random access of 10 blocks

$$\text{time} = 10(s + g_1 + b \times t)$$

$$= 251 \text{ ms.}$$

Sequential access time $\approx b \times ebt$ when b is very large

Random access time $\approx b \times (s + g_1 + b \times t)$

Heap File Organization

T_F = time to fetch a record

T_N = time to fetch next record

T_I = time to insert new record

T_U = time to update a record

T_D = time to delete a record

T_x = time for exhaustive reading of a file

T_y = time to reorganize a file.

b blocks: $T_F = \frac{b}{2} \times ebt$

Example: 100000 records of 400 bytes each

$$b = 16,667$$

$$T_F = 7 \text{ sec.}$$

fetch 10000 records

time ≈ 19 hours

$$T_N = T_F$$

Time to insert new record (T_I)

first the header will locate the block which will take time $s + r + btt$. Then the header will wait for the start to come.

$$\begin{aligned} T_I &= s + r + btt + (2r - btt) + btt \\ &= s + r + btt + 2r \end{aligned}$$

Time to delete a record (T_D)

$$T_D = T_F + 2r$$

to find the record and then to delete it.

$$T_x = b \times ebt \text{ (beginning to end)} = 14 \text{ sec.}$$

$$T_x \text{ (independent fetch of all records)} = n \times T_F = n \times \frac{b}{2} \times ebt$$

$$= 190 \text{ hours.}$$

~~T_F~~

$n \rightarrow$ no. of active records

$Bfr \rightarrow$ no. of records per block.

$$T_y = b \times ebt + \frac{n}{Bfr} \times ebt$$

Bucket

1 Bucket = n blocks, $n = 1, 2, 3, \dots$

1 Bucket = 2 blocks.

$$\text{time} = s + r + \underline{btt + ebt}$$

$$\text{Random time} = bk(s + r + dtt)$$

↓
no. of buckets

Value of dtt will change according to the number of blocks in a bucket.

$$\text{Random time} : bk(s + r + dtt) = bk(s + r) + bk \times dtt$$

If bucket size increases, then the no. of buckets bk falls but dtt rises $\therefore bk \times dtt$ remains constant but $bk(s + r)$ falls \therefore Total time falls.

Example: $bk = 1$ block = 2400 bytes

$$dtt = 0.8$$

Two block size = 2 blocks

$$dtt = 1.64$$

$$bk = 1250$$

~~$$1250 \times 0.84 = 1050$$~~

$$bk = 625$$

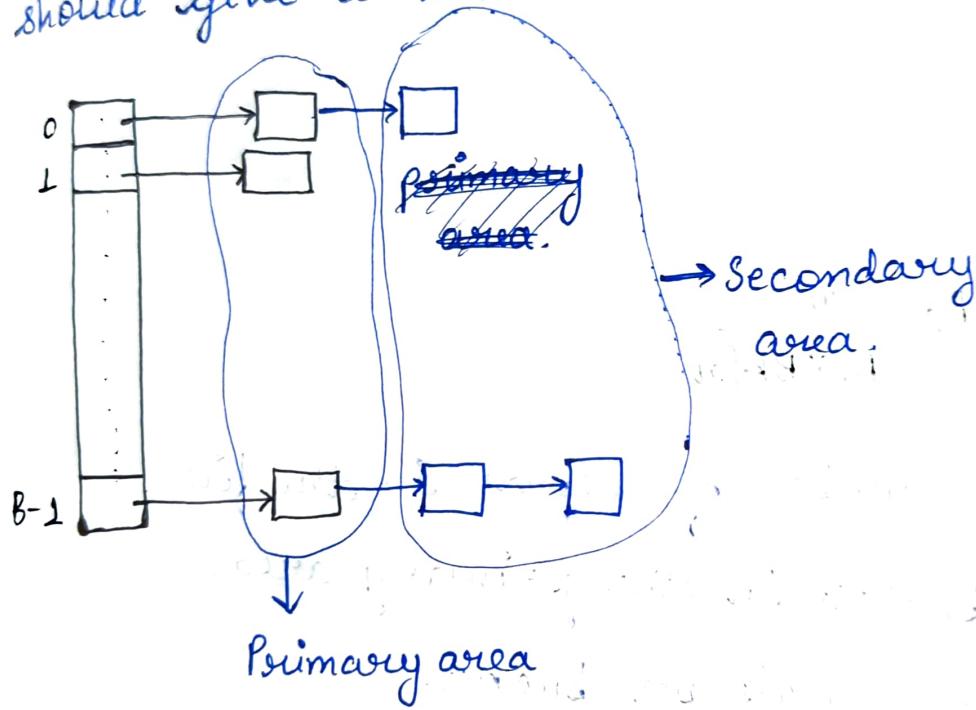
$$625 \times 1.64 = 1025$$

For random access, the time is reduced significantly but for sequential access, the reduction is marginal.

*Hash File Organization

Records are stored according to the key value. Before making the file, the number of records are estimated.

If the expected number of records is m , and the number of blocks reserved is n ($0 \leq n \leq m-1$). When a key value is supplied, the hash function should give a value between 0 and $n-1$.



When initially the file is formed, the records will be stored in the primary area and when the buckets are full, more records are stored in the secondary area.

If there is secondary area, then at max the whole chain at that particular position has to be searched.

Example : 22, 35, 33, 42, 71, 60, 47, 19, 46, 17

$$h(x) = x \bmod 5$$

No. of buckets = 5

No. of records per bucket = 3.

0	35	60	
---	----	----	--

1	71	46	
---	----	----	--

2	22	42	47	→	17		
---	----	----	----	---	----	--	--

3	33		
---	----	--	--

4	19		
---	----	--	--

Load factor : $L_f = \frac{n}{M \times B_{Kf}}$

$n \rightarrow$ expected number of records in the file

$M \rightarrow$ No. of buckets in the primary area.

$B_{Kf} \rightarrow$ No. of records per bucket.

$$\& T_f(\text{no chain}) = s+r+dtt$$

(i.e. when address of the bucket is known)

$$T_f(\text{successful}) = (s+r+dtt) + \frac{x}{2}(s+r+dtt)$$

In the chain, in secondary area, there are x number of buckets and in the primary area, one bucket.

When chain is created, records are allotted the slots which are free, \therefore they are not in a sequential order.

Time for fetching the first ~~record~~ ^{bucket} is $s + \alpha + dtt$ and then in ~~primary~~ secondary area, average time will be $\frac{x}{2}(s + \alpha + dtt)$.

$$T_f(\text{unsuccessful}) = s + \alpha + dtt + x(s + \alpha + dtt).$$

$$T_I = T_f(\text{unsuccessful}) + 2\alpha$$

First last block of the chain will be fetched and inserted in main memory but then it has to be overwritten, \therefore the head will have to wait for 2α .

$$T_D = T_f(\text{successful}) + 2\alpha$$

For deletion, it is considered, that the record is present in the file, \therefore first it is found and then deleted.

$T_U = T_D$; if key value is
(Updation) not changed

$T_D + T_I$; if key value is
changed.

For sequential organization, heap file organization is recommended.

* And for independent fetching of records, hash file organization is recommended.

$$T_x = n \times T_f$$

In this case, the buckets are almost independent more suitable for random access.

Extendible Hashing

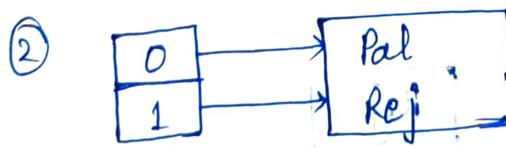
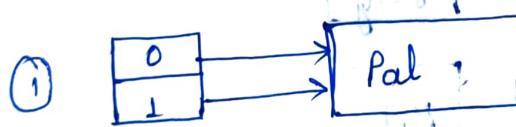
Initially no buckets are reserved. Buckets are created or removed according to need.

* Binary values of keys are considered. If the maximum number of bits is b, then i bits are considered for hashing.

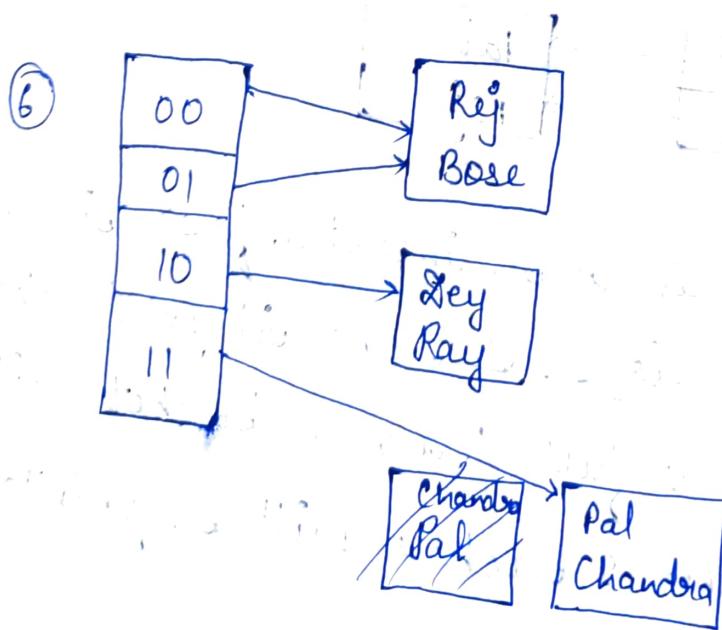
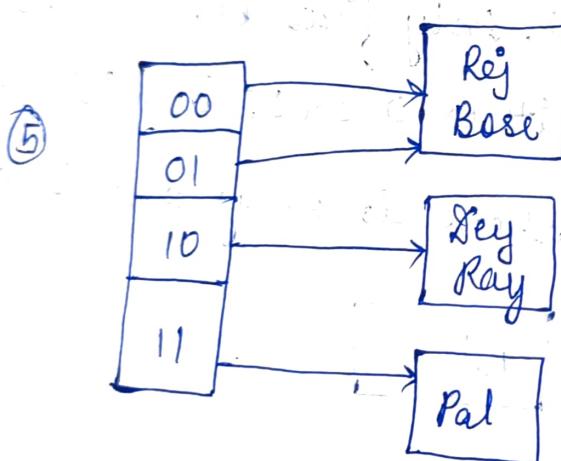
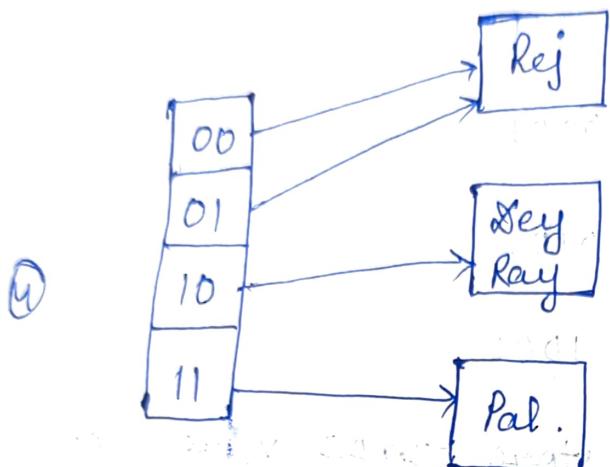
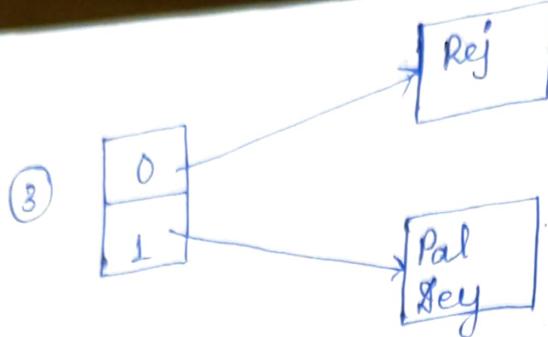
Name	hCName (hash values)
⑤ Rose	00101101
⑥ Chandra	11010101
③ Dey	10100011
⑦ Mathew	10000111
① Pal	11110001
④ Ray	10110101
② Rej	01011000

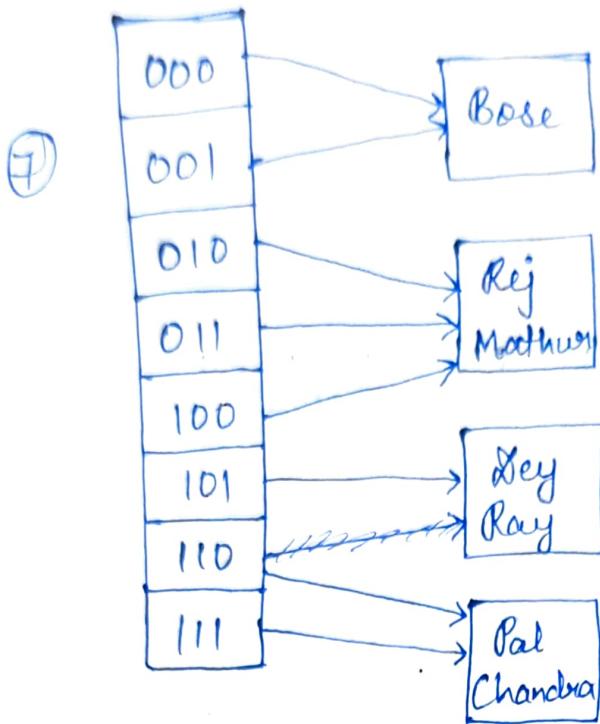
Each bucket can store atmost 2 records.

Let first record be Pal. Only first bit is considered.
 If the first bit is either 0, or 1; it will point to the same bucket as there is only one bucket.



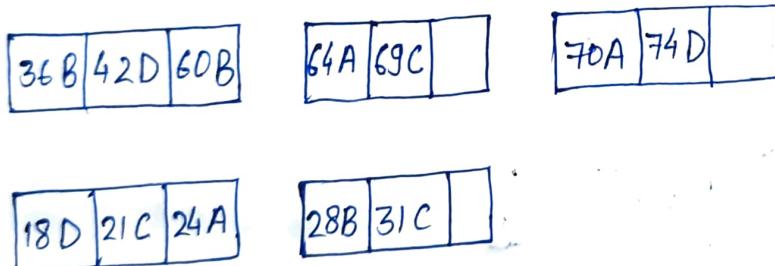
When Dey is to be inserted, first bit will be checked which is 1. But one bit points to a bucket which is full, ∴ a new bucket will be created. Now the 2 buckets will be pointed by 2 bits.



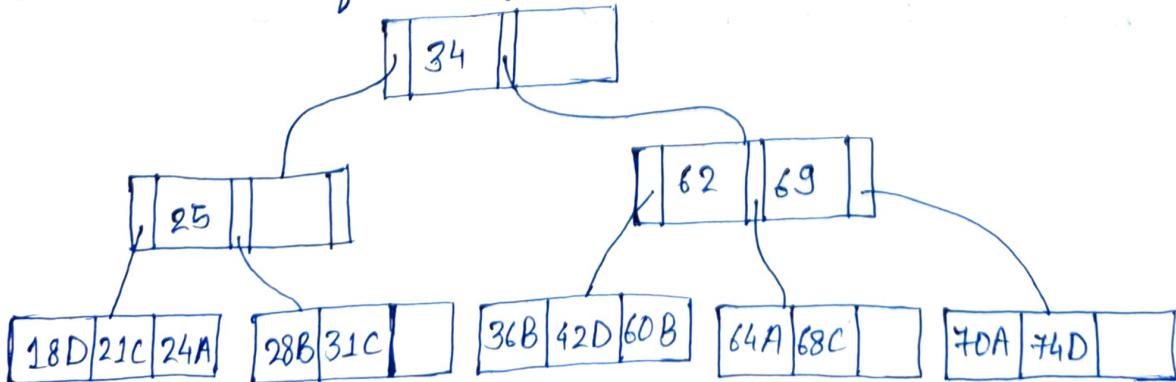


* B⁺ Trees

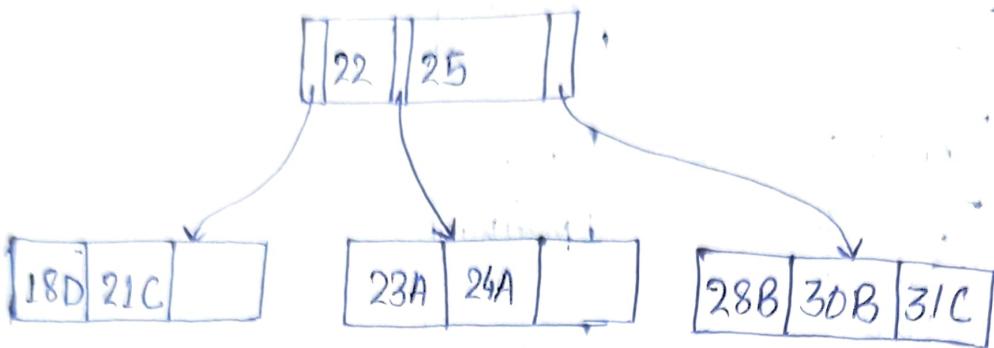
Let Roll No. and grade are being recorded. Let there be 2-3 records per leaf.



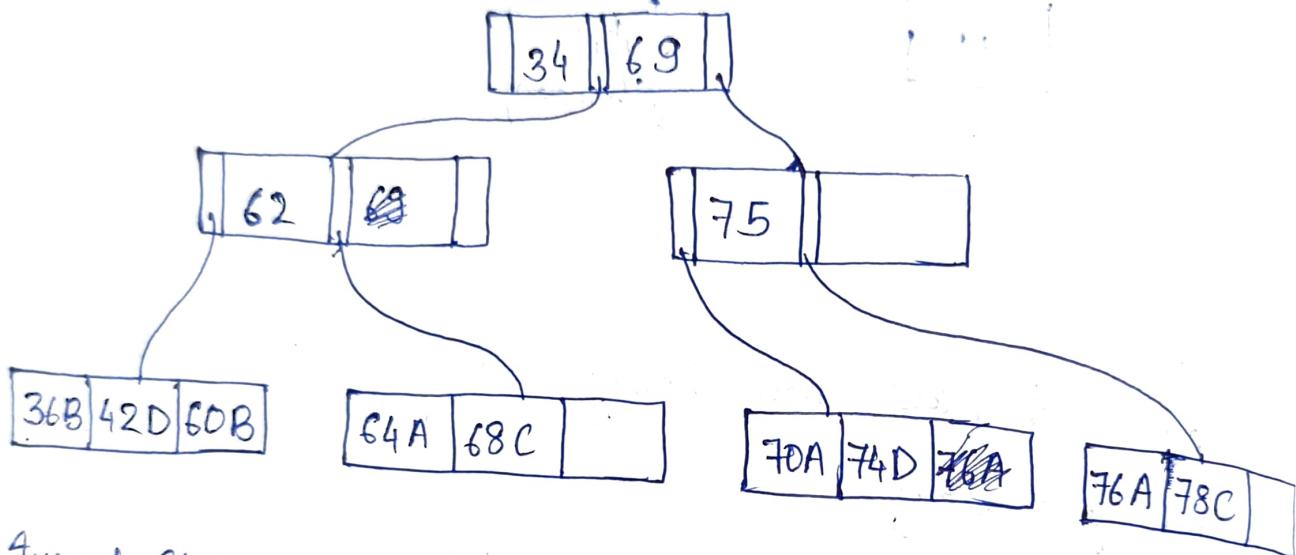
Leaves store data and non-leaves store pointers. It is not necessary that leaves and non-leaves have same number of data points and pointers.



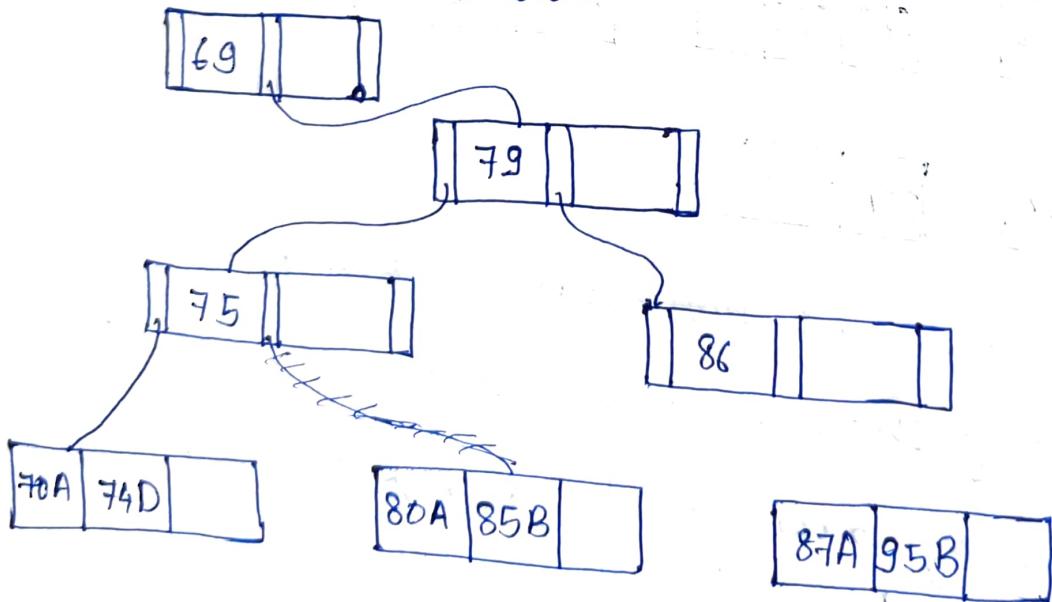
9) If 22A and 30B are inserted.



Insert 76A and 78C



Insert 80A, 85B, 87A and 95B.



14 A	14A		
------	-----	--	--

29 B	14A	29B	
------	-----	-----	--

11 C	11C	14A	29B
------	-----	-----	-----

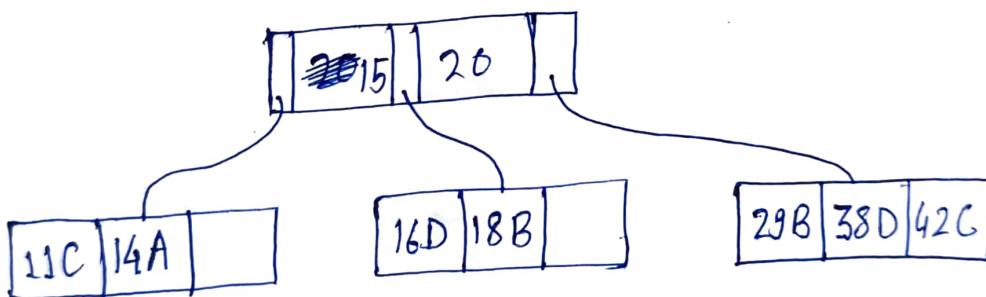
38 D	11C	14A	18B	29B	38D	
------	-----	-----	-----	-----	-----	--

Inserting 16 D,

20		
----	--	--

11C	14A	18B
-----	-----	-----

29B	38D	
-----	-----	--



det block size 800

$$\frac{800}{60} = 13 \text{ (No. of leaves)}$$

det record size 60

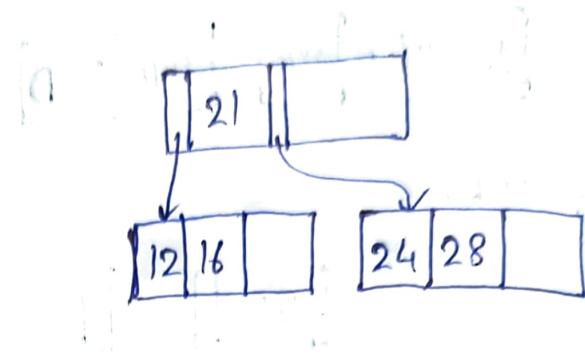
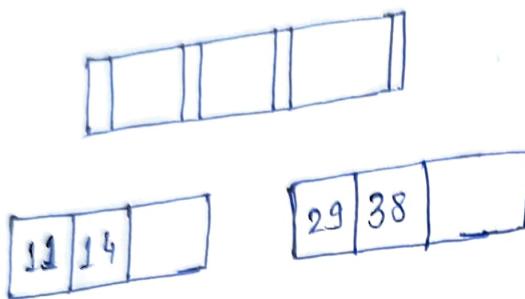
key size 20

pointer size 10

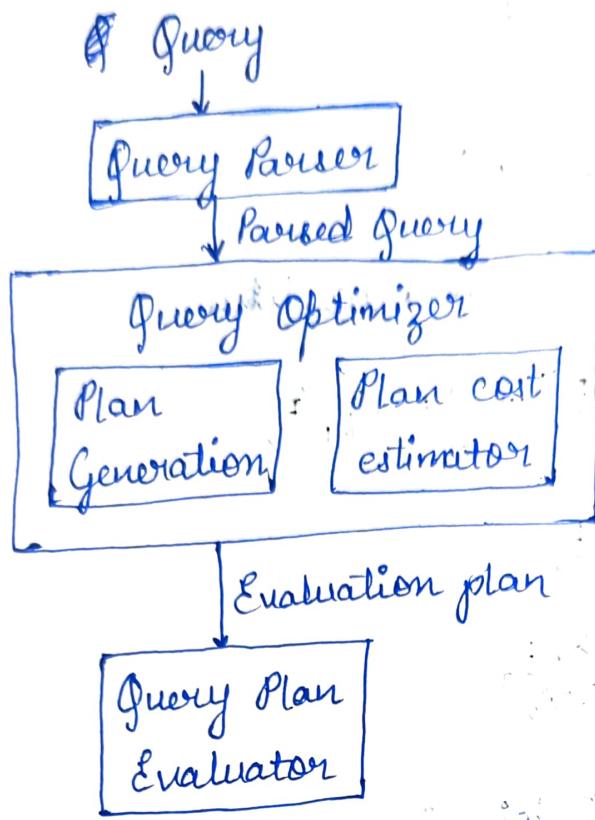
$$n \cdot 10 + (n-1) \cdot 20 \leq 800 \quad n = 27$$

2/3 records in a leaf except root
3/4/5 pointers in non-leaf except root

12, 16, 24, 28, 32, 20



Query Optimization



Query Parser: It checks whether the syntax of the query is correct or not and also checks if the attributes and relations are correct.
It converts the SQL query into relational algebra expression in the form of a tree.

Plan generation: It transforms the query into different equivalent forms.

Plan cost estimator: estimates the cost of each plan and chooses the plan with the optimum cost.

★ Equivalence Rules

$$1. \sigma_{\theta_1 \wedge \theta_2}(E) = \sigma_{\theta_1}(\sigma_{\theta_2}(E))$$

$$2. \sigma_{\theta_1}(\sigma_{\theta_2}(E)) = \sigma_{\theta_2}(\sigma_{\theta_1}(E))$$

Selection operation is commutative.

$$3. \pi_{L_1}(\pi_{L_2} \dots (\pi_{L_n}(E))) = \pi_{L_1}(E)$$

$$4. \sigma_{\theta}(E_1 \bowtie E_2) = E_1 \underset{\theta}{\bowtie} E_2$$

$$\sigma_{\theta_1}(E_1 \underset{\theta_2}{\bowtie} E_2) = E_1 \underset{\theta_1 \wedge \theta_2}{\bowtie} E_2$$

$$5. E_1 \underset{\theta}{\bowtie} E_2 = E_2 \underset{\theta}{\bowtie} E_1 \quad \theta \text{ join is commutative.}$$

$$6. (E_1 \underset{\theta_1}{\bowtie} E_2) \underset{\theta_2}{\bowtie} E_3 = E_1 \underset{\theta_1}{\bowtie} (E_2 \underset{\theta_2}{\bowtie} E_3)$$

Natural join is associative.

$$7. \sigma_{\theta_1}(E_1 \underset{\theta_2}{\bowtie} E_2) = (\sigma_{\theta_1}(E_1)) \underset{\theta_2}{\bowtie} E_2$$

$$\sigma_{\theta_1 \wedge \theta_2}(E_1 \underset{\theta_2}{\bowtie} E_2) = (\sigma_{\theta_1}(E_1)) \bowtie (\sigma_{\theta_2}(E_2))$$

when θ_1 involves only E_1 and θ_2 involves only E_2 .

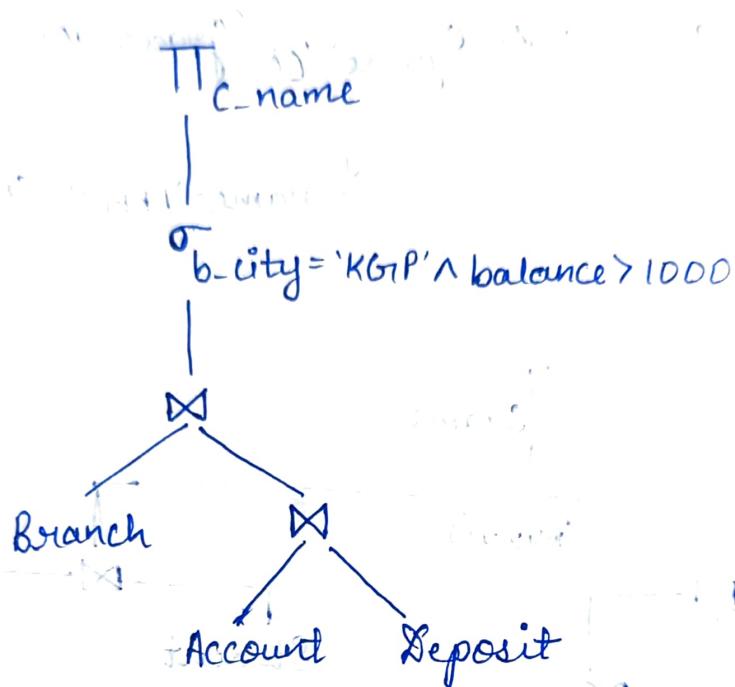
$$8. E_1 \cup E_2 = E_2 \cap E_1$$

Union and intersection are commutative and associative.

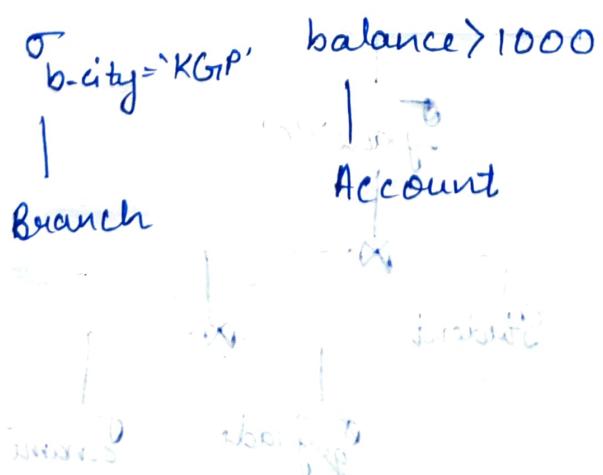
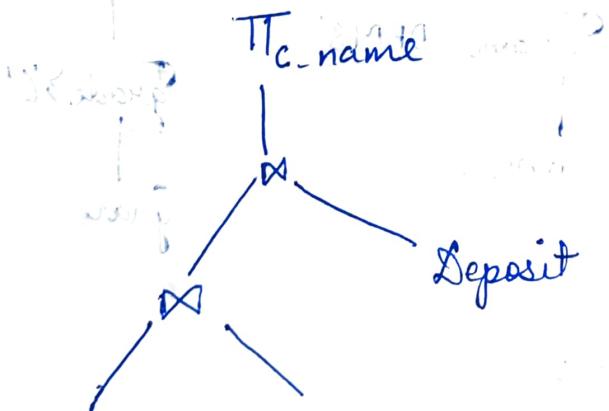
$$9. \sigma_{\theta}(E_1 - E_2) = \sigma_{\theta}(E_1) - \sigma_{\theta}(E_2)$$

$$10. \Pi_L(E_1 \cup E_2) = \Pi_L(E_1) \cup \Pi_L(E_2)$$

example: $\Pi_{C_name}(\sigma_{b_city='KGP' \wedge balance > 1000}(\text{Branch} \bowtie \text{Account} \bowtie \text{Deposit}))$



Join is moved up.

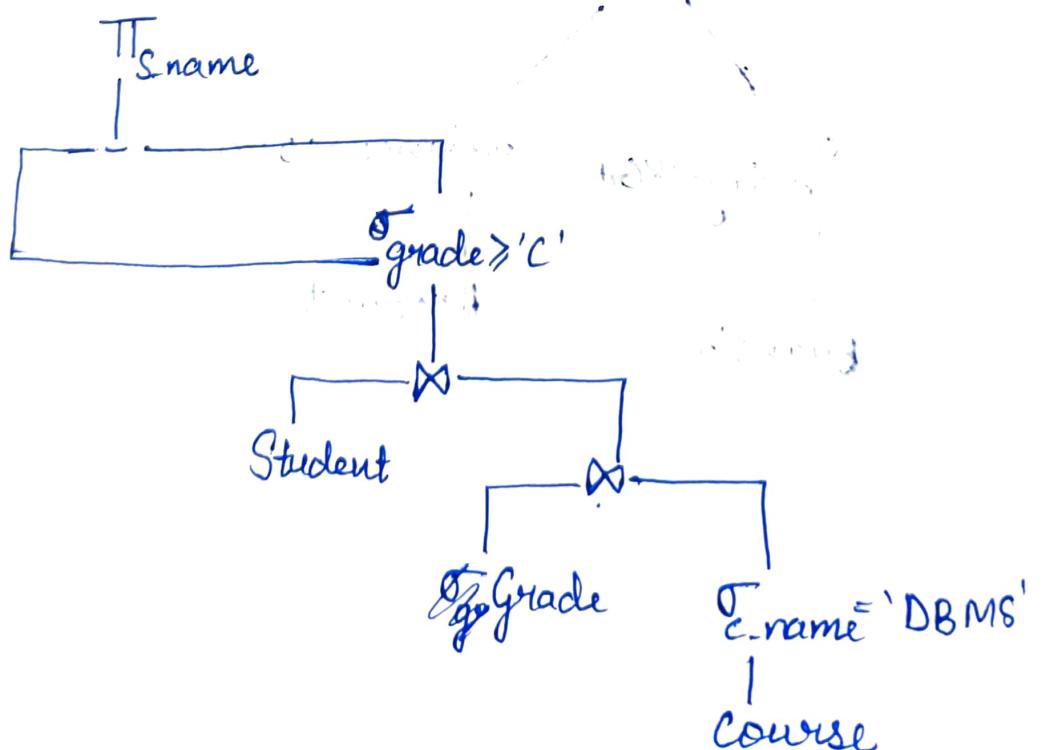
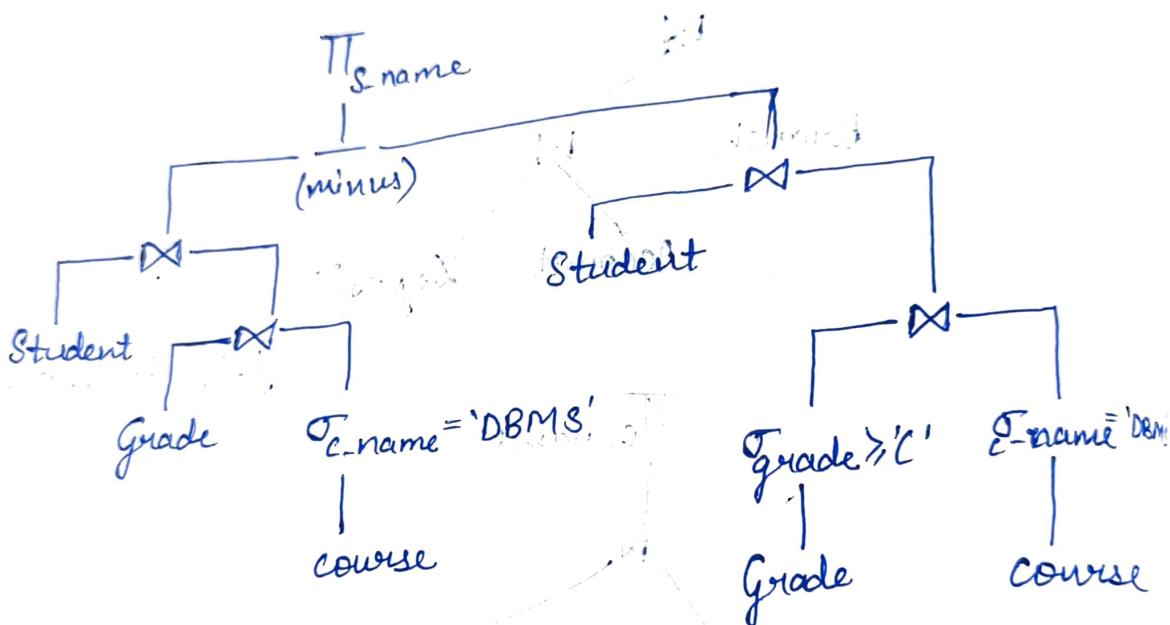


Example: Find the list of students who have not obtained grade 'C' or higher in 'DBMS' course.

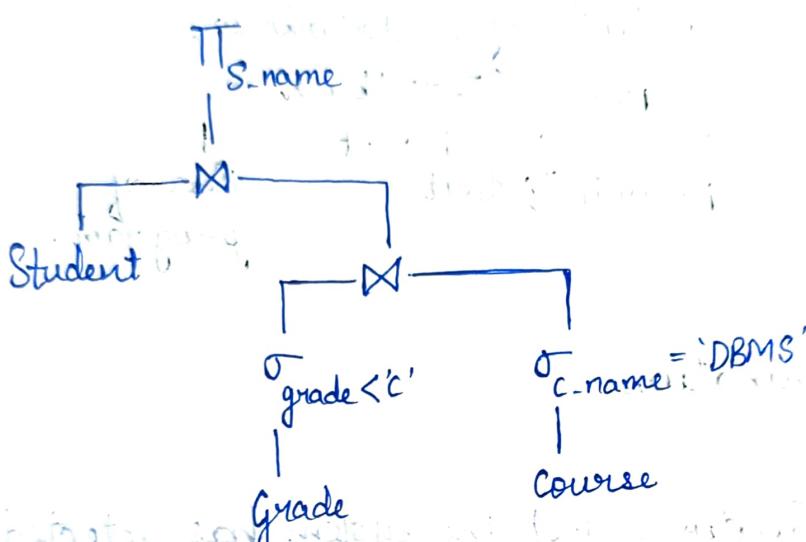
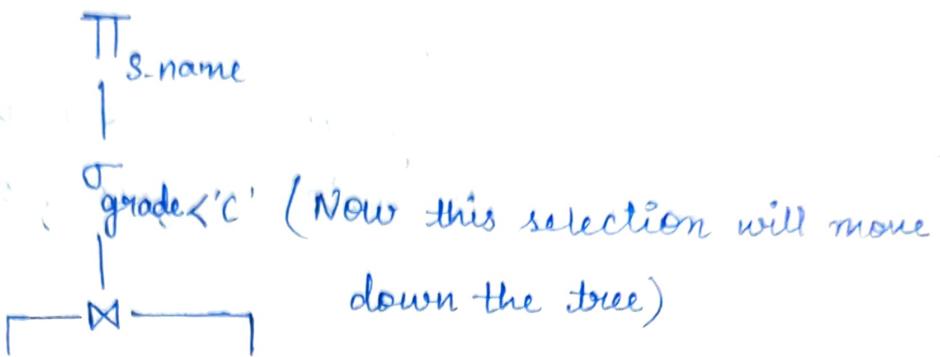
$\Pi_{S.name} (\text{Student} \bowtie (\text{Grade} \bowtie (\text{course}))) \sigma_{C.name = 'DBMS'} (course)$

- $(\text{Student} \bowtie (\sigma_{grade > 'C'} (\text{Grade})))$

$\sigma_{C.name = 'DBMS'} (course))$



$$f \circ \sigma_C(R) = \sigma_{\neg C}(R)$$



Recovery

~~if a & B~~ Backups have to be made in order to bring the data back to a previous stage at which data was correct in the event of failure.

Recovery Schemes

Forward Error Recovery

In this scheme when a particular error in the system is detected, the recovery system makes an accurate assessment of the state of the system and then makes appropriate adjustments based on anticipated

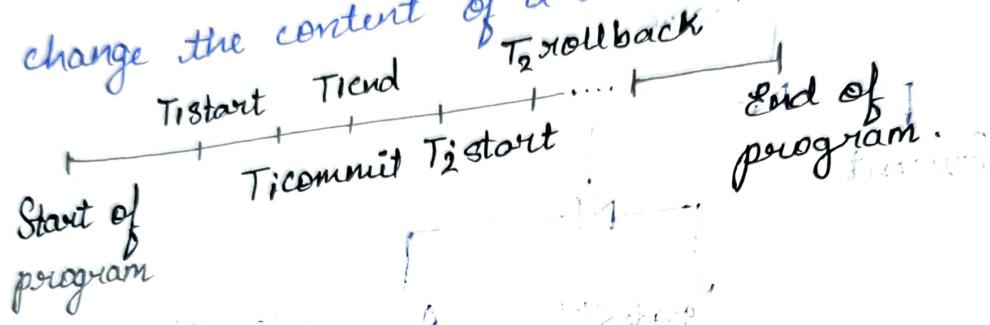
result had the system being error-free.

Backward Error Recovery

In this scheme the system is reset to some previous correct state which is free from errors.

• Transaction

Transaction is a problem unit whose execution may change the content of a database.



Successful termination

Suicidal termination: If the system has detected some error, and the system has rolled back that transaction

*Properties of a transaction

ACID Test

• Atomicity proper

This property of a transaction implies that it will run to completion as an indivisible unit at the end of which no changes have occurred to the database or the database has been changed in a consistent manner.

consistency property

This property implies that if the database was in a consistent state before the start of the transaction then on the termination of the transaction, the database will be in a consistent state.

isolation property

This property indicates that the actions performed by the transaction will be isolated or written outside of the transaction until the transaction is complete.

Durability property

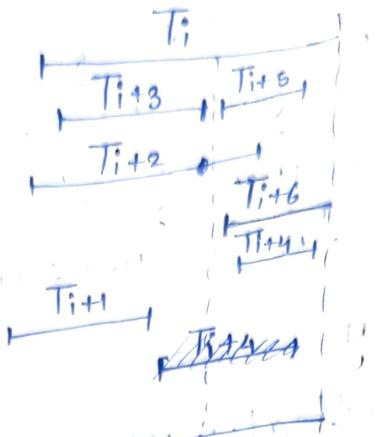
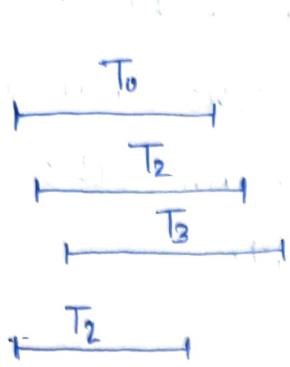
This property ensures that the commit action of a transaction will be reflected on the database on its termination.

*Recovery in Controlled DBMS

Log: It is a file in which all the transactions performed on the database is written. Trx start, commit and end markers are recorded.

*Write ahead log-strategy

It is a strategy to handle the transactions such that once the transaction has been started and modified some data and after committed all the changes done in the transaction will be recorded in the log file.



~~before~~ to check system points crossed when the system crashes only transactions ~~starts~~ ~~when the system crashes~~ ~~only transactions~~ ~~T_i and T_{i+6} are getting executed.~~

Checkpoint: Database is written to stable storage.

Transaction-consistent check point

Here the transactions that are active when system timer ~~signals~~ signals a check point are allowed to complete but no new transactions are allowed to be started until the check point has completed.

Action consistent check point

Active transactions are allowed to complete the current state before the checkpoint and no new actions can be started on the database until the checkpoint is complete.

transaction-oriented check point

take a checkpoint at the end of each transaction
by forcing the log of the transaction on stable
storage.

When the system crashes, transactions $T_{i+1}, T_{i+2}, T_{i+4}, T_{i+5}$,
and T_{i+6} transactions should be UNDO.

UNDO($T_i, T_{i+2}, T_{i+4}, T_{i+5}, T_{i+6}$)



REDO($T_{i+5}, T_{i+4}, T_{i+2}$)

UNDO(T_i, T_{i+6})

Do, undo, redo

Do operation

A transaction on the current database transforms
it from current state to new state. This is called
Do operation.

Undo operation comes into picture when a transaction
decides to terminate itself.

Transaction redo involves performing the changes made
by a transaction that committed before a system
crash. With a write-ahead log strategy, a committed
transaction implies that the log for the transaction would be
written to the secondary storage but the database

may or maynot have been identified before the system failure has been identified.

* Concurrency Management

* Schedule

The order in which the statement of transaction are executed is known as schedule.

T₁

1. Read (Aug-F-sal)
2. Aug-F-sal = f₁(Aug-F-sal)
3. write (Aug-F-sal)

S₁

1 2 3 4 5 6

S₂

1 4 2 5 3 6

T₃

1. Read (A)

2. A = A + 10

3. write(A)

T₂

4. Read (Aug-F-Salary)
5. Aug-S-sal = f₂(Aug-S-sal)
6. write (Aug-S-sal)

S₂

4 5 6 1 2 3

T₄

4. Read (A)

5. A = A * 1.1

6. write(A)

S₁

T₁ T₂

A = 231

S₂

T₂ T₃

A = 230

S₃

4 5 1 2 6 3

T₅

1. Read(A)
2. $A = A - 100$
3. write(A)
4. Read(B)
5. $B = B + 100$
6. write(B)

T₆

7. sum = 0
8. Read(A)
9. sum = sum + A
10. Read(B)
11. sum = sum + B
12. write(sum)

$$\begin{aligned}A &= 500 \\B &= 1100\end{aligned}$$

1 7 8 2 3 9 4 5 6 10 11 12

Inconsistent read problem: when a system allows to read a variable 2 times, then the 2nd read statement might be inconsistent as it does not read the value from the first read.

Given a set of transactions, the schedule will be serialisable if following conditions hold.

(i) all transactions are correct in the sense that if any one of the transactions is executed by itself on a consistent database, the result ~~must~~ databases will be consistent.

(ii) any serial execution of the transaction is also correct & preserve the consistency of the database. The result obtained is also correct.

a directed edge from node T_i to T_j indicates one of the following:

1. T_i performs the operation $\text{Read}(A)$ & then T_j performs the operation $\text{write}(A)$.
~~Read(A)~~

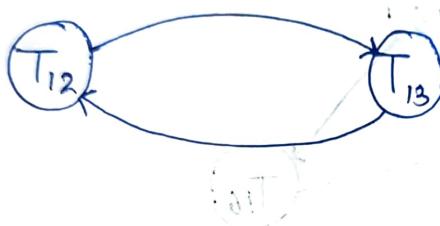
2. T_j performs the operation $\text{write}(A)$ then T_i performs the operation $\text{read}(A)$.

<u>Schedule</u>	<u>T_9</u>	<u>T_{10}</u>	<u>T_{11}</u>
Read(A)	✓		
$A = f_1(A)$	✓		
write(A)	✓		
Read(A)			✓
$A = f_2(A)$			✓
write(A)		✓	
Read(B)		✓	
$B = f_3(B)$		✓	
write(B)		✓	
Read(B)			✓
$B = f_4(B)$			✓
write(B)			✓



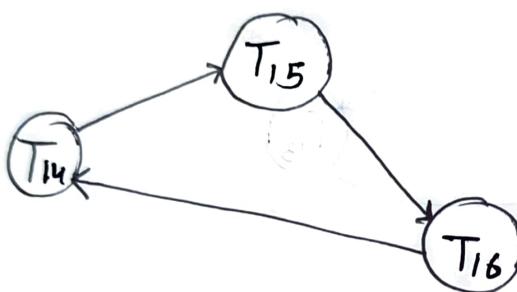
~~for~~ Schedule

	<u>T₁₂</u>	<u>T₁₃</u>
Read(A)	✓	
$A = f_1(A)$	✓	
Read(A)		✓
$A = f_2(A)$		✓
write(A)		✓
write(A)	✓	



	<u>S</u>	<u>T₁₄</u>	<u>T₁₅</u>	<u>T₁₆</u>
Read(A)		✓		
Read(B)			✓	
$A = f_1(A)$		✓		
Read(C)				✓
$B = f_2(B)$			✓	
write(B)			✓	
$C = f_3(C)$				✓
write(C)				✓
write(A)		✓		
Read(B)				✓
Read(A)			✓	

<u>S</u>	<u>T₁₄</u>	<u>T₁₅</u>	<u>T₁₆</u>
$A = f_4(A)$		✓	
Read(C)	✓		
write(A)		✓	
$C = f_5(C)$	✓		
write(C)	✓		
$B = f_6(B)$			✓
write(B)			✓



A Bad Locking Scheme

• **LOCK :** Lock is a variable associated with various data items.

• **Exclusive lock**: Whenever locked, transaction can write or update while no other transaction can even read (exclusive use to one item only).

• **Shared lock**:

One can read data item & can work with it but ~~can't modify it~~. Other transactions can also read.

T₅

T₆

lock+(sum)

✓

sum = 0

✓

locks(A)

✓

Read(A)

✓

sum = sum + A

✓

unlock(A)

✓

lockxc(A)

✓

Read(A)

✓

A = A - 100

✓

write(A)

✓

unlock(A)

✓

lockxc(B)

✓

Read(B)

✓

B = B + 100

✓

write(B)

✓

unlock(B)

✓

lockxc(B)

✓

Read(B)	✓
sum = sum + B	✓
write(sum)	✓
unlock(B)	✓
unlock(sum)	✓

• Two Phase Locking

It has 2 phases, a growing phase where the number of locks increases from 0 to a maximum value for the transaction and one contracting phase where the no. of locks held decreases from max. value to zero.

If a schedule follows 2 phase locking then the series is serializable.

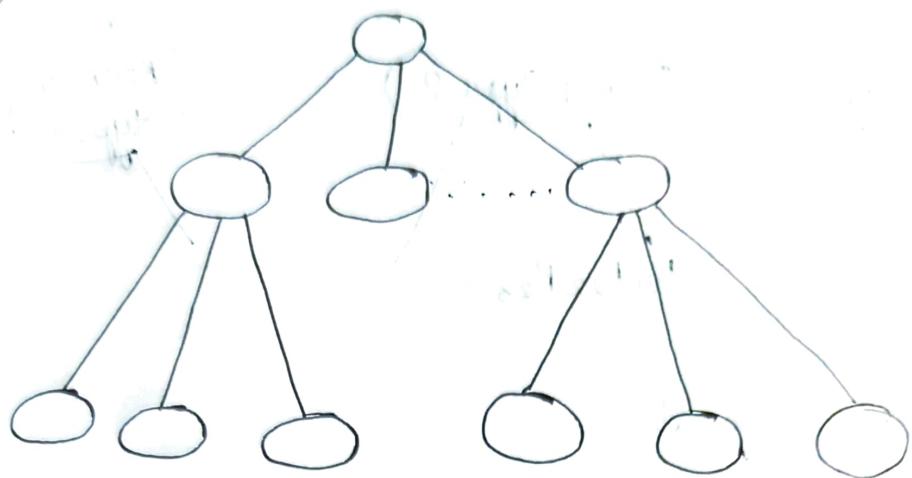
Let there be T_1, T_2, \dots, T_x transactions and assume there is a cycle.

$$T_a \rightarrow T_b \rightarrow T_c \rightarrow \dots \rightarrow T_x \rightarrow T_a$$

T_b 's lock will be executed ~~after~~ ^{after} T_a 's unlock is performed. But after T_x , T_a cannot again lock the data items according to 2-phase locking.

Any data item in a database can be locked only once

Intention Share (IS)
When a ~~node~~ node is locked in IS mode, then the node or its descendants cannot be exclusively locked. The descendants may be locked in S or IS mode.

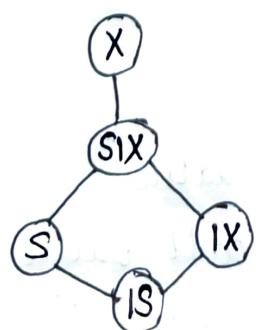


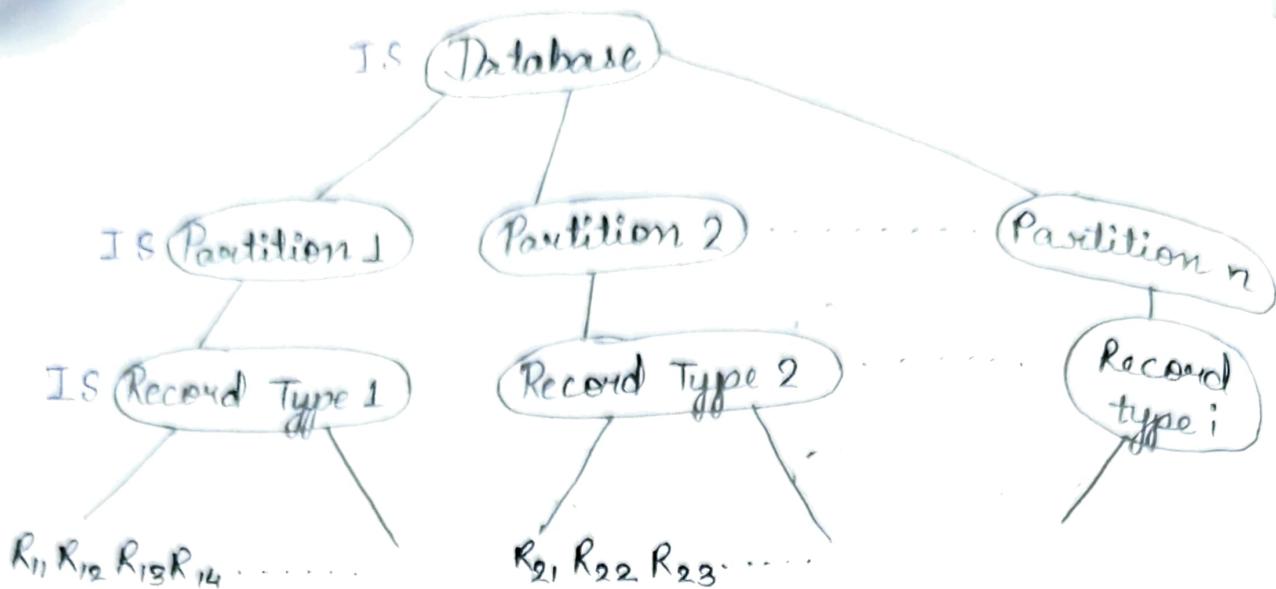
Intention Exclusive (IX)

When a node is locked in IX mode then the node itself ~~be~~ cannot be exclusively locked. However, the descendants can be locked in any of the locking modes.

Shared and Intention Exclusive (SIX)

The node locked in SIX mode and all the descendants ~~and all the descendants~~ will intrinsically locked in shared mode. Any of the descendants can be locked in X or SIX mode.





$R_{13} \rightarrow$ To be retrieved

$R_{13} \rightarrow S$

$R_{22} \rightarrow$ update $R_{22} \rightarrow X$

Time Stamp Based Ordering Protocol

When a transaction starts executing then it is assigned some value which is equal to the system clock value when the transaction started.

$T_i \rightarrow$ time stamp value t_i

$T_i(t_i) \& T_j(t_j)$

s.t. $t_i < t_j$

$T_i \rightarrow$ older transaction $T_j \rightarrow$ younger transaction

$X: \{x, wx, Rx\}$

$x:$ value of X

$wx:$ write timestamp value

~~Rx~~ It is the system call value when the data item was written for the last time.

R_x: Read timestamp value.

The system ~~will~~ clock value when the transaction last read the data item.

• Read operation

A transaction T_a with timestamp value t_a issues a ~~read~~ ^{read} operation for the data item $X \{x, w_x, R_x\}$. ~~The~~ op will

1) This request will be allowed if $t_a > w_x$.

After the op : $X \{x, w_x, t_a\}$

2) This request will fail if $t_a < w_x$.

T_a is an older transaction than the one which wrote the value of X .

• Write operation

A transaction T_a with timestamp value t_a issues a write operation for data item $X \{x, w_x, R_x\}$.

1) $t_a > w_x$ and $t_a > R_x$ then the write operation will be allowed. The transaction which ~~has~~ last read and wrote the value of X should be older than T_a . w_x will be replaced by t_a .

2) $t_a < R_x$: A longer transaction is already using the value of X : T_a is not allowed to modify the value of X . T_a is rolled back and a new timestamp value is provided

3) $R_x < t_a < w_x$, younger transaction has already written the

value of x , and T_0 is not allowed to write the value of x .

<u>Schedule</u>	$T_{22}(t_{22})$	$T_{23}(t_{23})$
1. $\text{Sum} = 0$	✓	
2. $\text{Read}(A)$	✓	
3. $\text{Sum} = \text{Sum} + A$	✓	✓
4. $\text{Sum} = 0$		✓
5. $\text{Read}(A)$		✓
6. $A = A - 100$		✓
7. $\text{Write}(A)$		
8. $\text{Read}(B)$	✓	
9. $\text{Sum} = \text{Sum} + B$	✓	
10. $\text{Show}(\text{Sum})$	✓	
11. $\text{Sum} = \text{Sum} + A$		✓
12. $\text{Read}(B)$		✓
13. $B = B + 100$		✓
14. $\text{Write}(B)$		✓
15. $\text{Sum} = \text{Sum} + B$		✓
16. $\text{Write}(\text{Sum})$		✓
Show		

Initially, $A: 400, w_a, R_a$ $B: 500, w_b, R_b$

Step(2) : $A: 400, w_a, t_{22}$ "

Step 5 : $A: 400, w_a, t_{23}$ "

Step 7 : $A: 300, t_{23}, t_{23}$ "

Step 8 :

"

B: 500, w_b, t₂₂

Step 10:

sum = 900

Step 12 :

"

B: 500, w_b, t₂₃

Step 14 :

"

B: 600, t₂₃, t₂₃

Schedule

T₂₂

T₂₃

1) sum = 0		✓
2) Read(A)		✓
3) A = A - 100		✓
4) write(A)		✓
5) Read(A)	✓	
6) sum = sum + A		✓
7) Read(B)		✓
8) B = B + 100		✓
9) write(B)		✓
10) sum = sum + B		✓
11) show(sum)		✓
12) sum = 0	✓	$\nexists T_{22} : t'_{22} (\geq t_{23})$
13) Read(A)	✓	
14) sum = sum + A	✓	
15) Read(B)	✓	
16) sum = sum + B	✓	
17) show(sum)	✓	

A: 400, w_b, R_a

B: 500, w_b, R_b

Step 2 A: 400, w_a, t₂₃

"

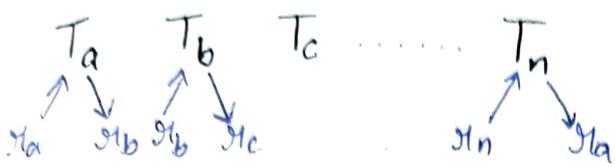
Step 3 A: 300, t₂₃, t₂₃

"

Step 4

Step 5 Read ~~w~~ is not allowed therefore it will rollback and will be forced to be executed in the end serially.

Deadlock and its resolution



Wait-for graph

T_id	Items-locked	Items-waiting-for
------	--------------	-------------------

T ₂₇	B	C, A
-----------------	---	------

T ₂₈	C, M	H, G
-----------------	------	------

T ₂₉	H	D, E
-----------------	---	------

T ₃₀	G	A
-----------------	---	---

T ₃₁	A, E	<u>C*</u> (added later)
-----------------	------	-------------------------

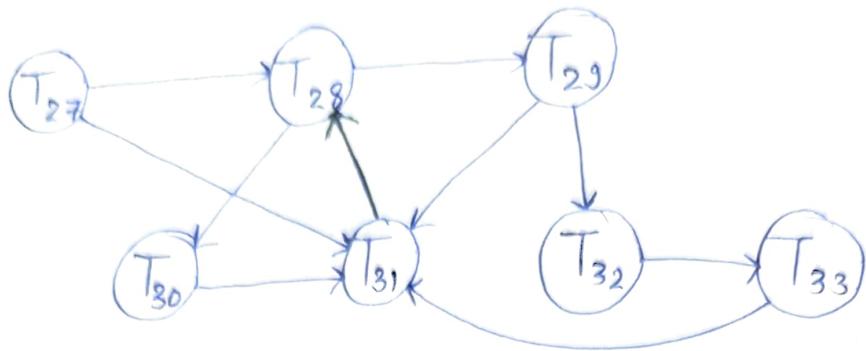
T ₃₂	D, I	F
-----------------	------	---

T ₃₃	F	E
-----------------	---	---

T_i is waiting for a data item currently allocated to and is being used by T_j , then there will be a directed edge from T_i to T_j .



If there is a cycle in the wait-for, then there is a deadlock.



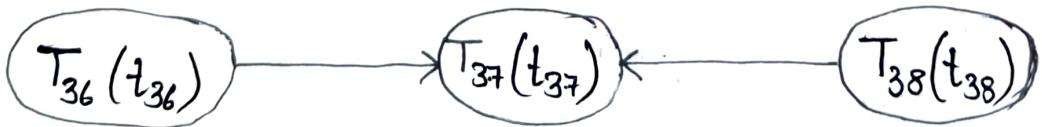
* After adding c^* in the waiting list of T_3 , then a cycle is created.

* How to prevent deadlock?

* Wait-Die Scheme

Wait and die scheme for deadlock prevention is as follows:

- If the requesting trnx is older than the trnx that holds the lock on the requested data-item, the requesting trnx is allowed to wait.
- If the requesting trnx is younger than the trnx that holds the lock on the requested data-item, the requesting trnx is aborted and rolled back.



$$t_{36} < t_{37} < t_{38}$$

$T_{36} \rightarrow \text{wait}$, $T_{38} \rightarrow \text{abort (Roll-back)}$

* Abandon - Wait Scheme

- i) If a younger transaction holds a data item requested by an older one, younger transaction will be aborted and rolled back.
- ii) If a younger transaction requests a data item held by older transaction, the younger transaction is allowed to wait.



$$t_{39} < t_{40}$$

$t_{40} \rightarrow$ aborted and rolled back.

* Distributed Database (DDB) {DDBMS}

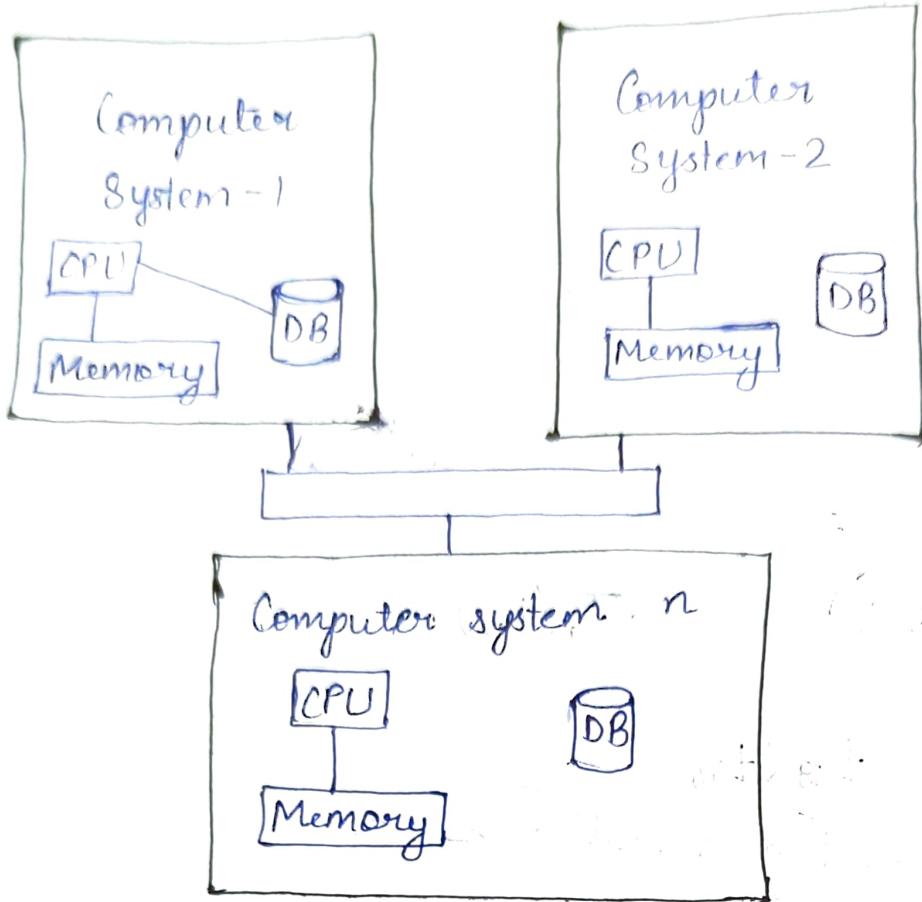
Database is stored at multiple locations

- Shared Memory architecture

Multiple processors share the secondary disk as well as the primary memory also.

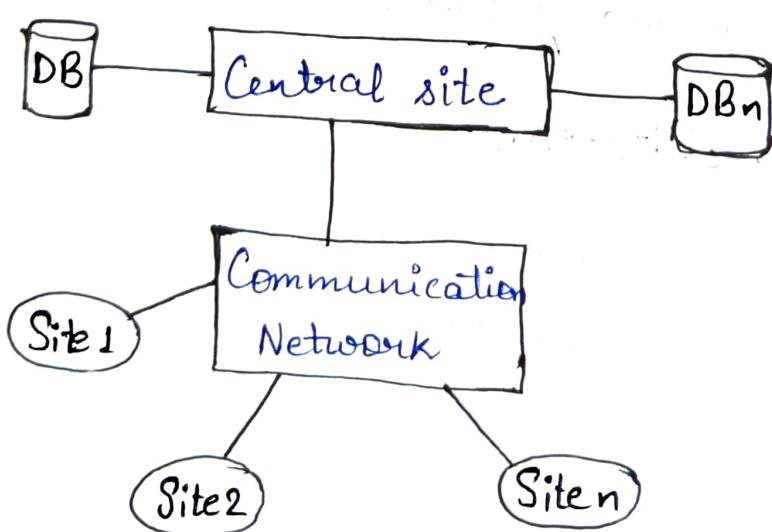
- Shared Disk architecture

Multiple processors share the same secondary disk but they have their own memory.



In the above architecture, all databases have their own primary and secondary memory but they are connected by network to share information.

- **Centralized database.**



Distributed DB



Databases are distributed on different sites. When a query comes it is directed accordingly.

* Advantages

- Management of distributed Data with different levels of transparency

Distribution transparent

the ~~physical to~~ physical (memory) location of the file will not be known by the user.

Network transparency

This refers to freedom for the user from the operational details of the network.

Location transparency: When we want to execute a command, it ~~to~~ should be independent of the location of the database.

Naming transparency: Once a name is specified, the name object can be accessed unambiguously without ~~addit~~ additional information.

- Replication transparency

It makes the user unaware of the existence of multiple copies.

Better availability: If some site is down, then data can be used from a copy in another site.

Performance *

Reliability

- Fragmentation transparency

Horizontal fragmentation: Database is fragmented row wise or record-wise.

Vertical fragmentation: Column-wise.

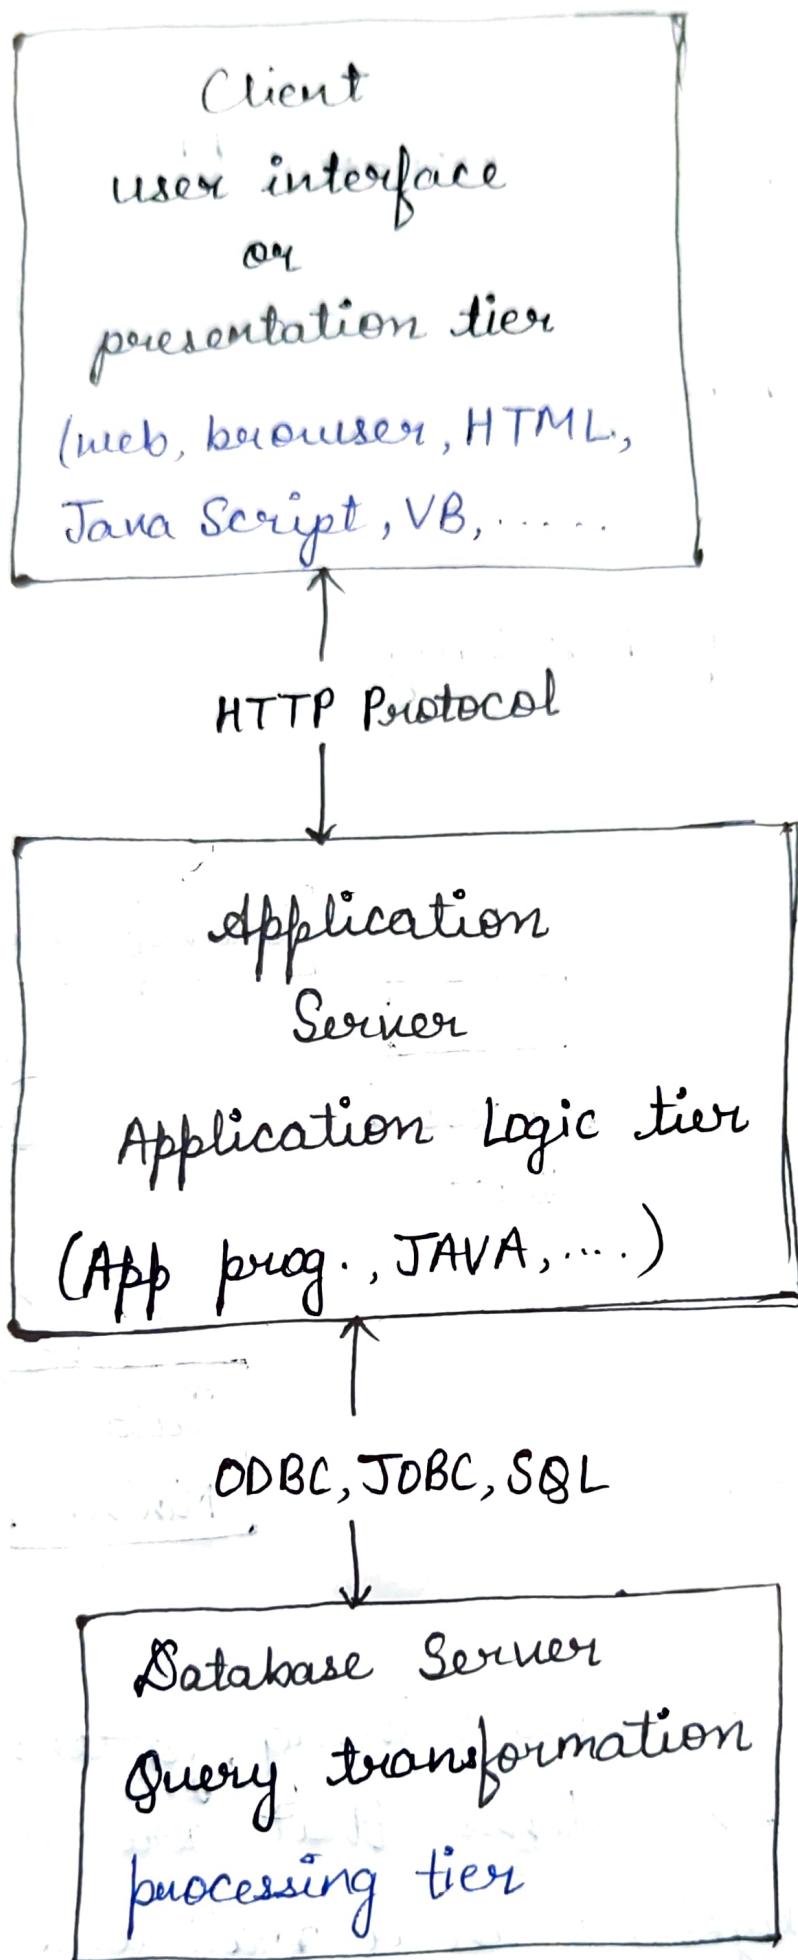
Mixed fragmentation: Both horizontal and vertical fragmentation.

- Increased reliability and availability

~~Reliability~~ Reliability is the property that a system is not down at a given time.

- Improved performance

- Easier expansion



* Data Mining

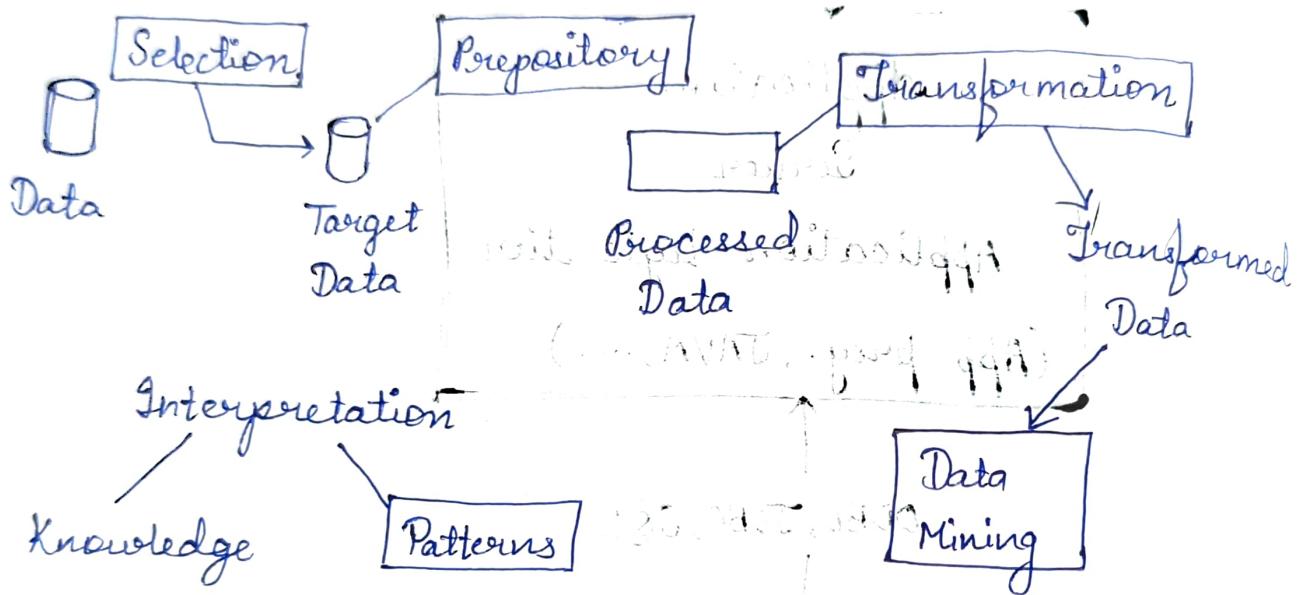
Knowledge Discovery in Databases (KDD)

In the ~~KDD~~ KDD process, there are 3 major parts:

* Association Rule Mining

* Classification

* Clustering / Similarity Matching



In case of association rule mining, patterns are formed such that the items that are more likely to be purchased together in a mall are kept together based on history of purchases.

In case of classification, when new data is acquired using the previous classification how can new predictions be made in order to increase profit.
Example: New loans to be offered to new potential

customers.
In case of clustering, data points are ~~clarify~~ assigned to different clusters.

TID	Item
1	Bread, Coke, Milk
2	Beer, Bread
3	Beer, Coke, Diaper, Milk
4	Beer, Bread, Diaper, Milk
5	Coke, Diaper, Milk

Item Set

Item set is a collection of one or more items.

K-item set

An item set which contains k items.

Support Count (σ)

Frequency of occurrence of an item set in a set of transactions.

$$\sigma \{ \text{Milk, Beer, Diaper} \} = 2$$

Support

Fraction of transactions that contain an item set.

$$S \{ \text{Milk, Beer, Diaper} \} = \frac{2}{5}$$

Frequent Item Set

It is an item set whose support > Minimum support threshold.

Association Rule discovery

$\{Diaper\} \rightarrow \{Beer\}$

$\{Milk, Bread\} \rightarrow \{Coke\}$.

→ Person who purchases milk and bread is also likely to purchase Coke.

Let there be d unique items

$$\text{Total no. of itemset} = 2^d$$

$$\text{Total no. of possible association rules} = 3^d - 2^{d+1} + 1$$

* Apriori Algorithm

1. Calculate Minimum support.
2. Find out frequent itemset. (Any subset of frequent itemset must be frequent).
∴ If a set is not frequent then any of its superset will also not be frequent.

L_k : Set of candidate k -itemset

frequent

candidate

Join step:

L_k is generated by joining L_{k-1} with itself where L_k is the candidate itemset of size, L_{k-1} is the frequent itemset of size k .

Prune step: $(k-1)$
any $k-1$ itemset that is not frequent cannot be a
subset of frequent k itemset.
 $L_1 : \{\text{frequent items}\};$

for ($k=1$; $L_k = \emptyset$; $k++$) do begin

C_{k+1} = candidates generated from L_k ;

for each transaction t in database

DO

increment the count of all
candidates in C_{k+1} that are
contained in t

L_{k+1} = candidates in C_k which are frequent
according to support.

TID	Items	
T1	I1, I2, I5	
T2	I2, I4	
T3	I2, I3	$\min \text{ support} =$
T4	I1, I2, I4	$\frac{2}{9} = 22\%$
T5	I1, I3	
T6	I2, I3	Minimum
T7	I1, I3	confidence = 70%.
T8	I1, I2, I3, I5	
T9	I1, I2, I3	

Final step: ($k-1$)
 Any $k-1$ itemset that is not frequent cannot be a
 subset of frequent k itemset.
 L_k : {frequent items};

for ($k=1$; $L_k = \emptyset$; $k++$) do begin
 C_{k+1} = candidates generated from L_k ;
 for each transaction t in database
DO

increment the count of all
 candidates in C_{k+1} that are
 contained in t

L_{k+1} = candidates in C_k which are frequent
 according to support.

TID	Items	
T1	I1, I2, I5	
T2	I2, I4	
T3	I2, I3	
T4	I1, I2, I4	
T5	I1, I3	
T6	I2, I3	
T7	I1, I3	
T8	I1, I2, I3, I5	
T9	I1, I2, I3	

min support = $\frac{2}{9} = 22\%$

Minimum confidence = 70%.

(81)

IT	SC
$\{I_1\}$	6
$\{I_2\}$	7
$\{I_3\}$	6
$\{I_4\}$	2
$\{I_5\}$	2

C₁

IT	SC
$\{I_1\}$	6
$\{I_2\}$	7
$\{I_3\}$	6
$\{I_4\}$	2
$\{I_5\}$	2

L₁

(82)

L₁ Join L₁

$\{I_1, I_2\}$	4
$\{I_1, I_3\}$	4
$\{I_1, I_4\}$	1
$\{I_1, I_5\}$	2
$\{I_2, I_3\}$	4
$\{I_2, I_4\}$	2
$\{I_2, I_5\}$	2
$\{I_3, I_4\}$	0
$\{I_3, I_5\}$	1
$\{I_4, I_5\}$	0

C₂

$\{I_1, I_2\}$	4
$\{I_1, I_3\}$	4
$\{I_1, I_5\}$	2
$\{I_2, I_3\}$	4
$\{I_2, I_4\}$	2
$\{I_2, I_5\}$	2

L₂

(83)

L₂ Join L₂

$\{I_1, I_2, I_3\}$	2
$\{I_1, I_2, I_5\}$	2

C₃

$\{I_1, I_2, I_3\}$	2
$\{I_1, I_2, I_5\}$	2

L₃

$\{I_1, I_2, I_3\}$
hypersets of non-frequent sets of 2 items will not
be frequent, \therefore we will not consider them.

③ $C_4 : L_3 \text{ Join } L_3$
 $C_4 \neq \emptyset$

$\{I_1, I_2, I_3, I_5\}$ is not frequent because $\{I_3, I_5\}$ is not
frequent.

$$L = L_1 \cup L_2 \cup L_3$$

$$L^* = \{\{I_1\}, \{I_2\}, \{I_3\}, \{I_4\}, \{I_5\}, \{I_1, I_2\}, \{I_1, I_3\}, \{I_1, I_5\}, \\ \{I_2, I_3\}, \{I_2, I_4\}, \{I_2, I_5\}, \{I_1, I_2, I_5\}, \{I_1, I_2, I_5\}\}$$

Association Rule:

I, s is subset of I

$s \rightarrow (I - s)$ if

$$\frac{\text{support-count}(I)}{\text{support-count}(s)} > \text{min. confidence}$$

$$\{I_1, I_2, I_5\}$$

$$\{I_1, I_2\}, \{I_1, I_5\}, \{I_2, I_5\}, \{I_1\}, \{I_2\}, \{I_5\}$$

$$R1: I_1, I_2 \rightarrow I_5 \quad \frac{SC\{I_1, I_2, I_5\}}{SC\{I_1, I_2\}} = 50\%$$

Reject.

$R_2: I_1, I_5 \rightarrow I_2$

R_2 is selected

$$\frac{SC\{I_1, I_2, I_5\}}{SC\{I_1, I_5\}} = 100\%,$$